

# Policy Methods and PPO

Matteo Cozzi

April 15, 2025



# A Step Back: What is our Goal in RL?

## Reinforcement Learning (RL)

Reinforcement Learning is a framework where an agent learns to make decisions through trial and error, by interacting with an environment and receiving feedback in the form of rewards.

- The agent takes actions in an environment.
- The environment responds with a new situation and a reward signal.
- Over time, the agent aims to learn a strategy (policy) that chooses actions to maximize long-term reward.

## Core Objective

Learn a policy that leads to the most rewarding behavior over time.

# A Step Back: Value-based vs Policy-based Methods

## Value-based Methods

- Learn value function  $Q(s, a)$ .
- The policy is *indirectly* derived from the learned value function:

$$\pi(s) = \arg \max_{a \in \mathcal{A}(s)} Q(s, a).$$

- Examples: **Q-Learning**, **Deep Q-Network (DQN)**.

## Policy-based Methods

- Directly parameterize the policy  $\pi(a | s; \theta)$  and optimize it with respect to expected returns.
- Examples: **REINFORCE**, **Actor-Critic**, **PPO**.

# Why Move to Policy-Based Methods?

- **Direct Optimization:**

- Policy-based methods directly optimize the decision-making strategy.
- This often leads to more stable convergence in complex environments.

- **Handling Continuous Actions:**

- Unlike value-based methods, policy-based methods naturally extend to continuous or high-dimensional action spaces.
- No need to perform maximization over actions (e.g.,  $\max_a Q(s, a)$ ), which may be intractable.

- **Stochastic Policies:**

- Can represent and learn stochastic behaviors, useful in partially observable or multi-modal settings.

- **Foundation for Advanced Methods:**

- Forms the basis for powerful algorithms like Actor-Critic and Proximal Policy Optimization (PPO).

# From Policy-Based to Gradient-Based Optimization

- **Policy-based methods** directly parameterize a policy  $\pi_{\theta}(a | s)$  and optimize a performance measure  $J(\theta)$ .
- Several optimization approaches exist:
  - **Gradient-based methods:** Optimize using the gradient  $\nabla_{\theta} J(\theta)$ . This is the most common approach in deep reinforcement learning.
  - **Gradient-free methods:** Such as evolutionary strategies or other black-box optimizers.
- In modern deep RL, **policy gradient methods** are predominantly used due to their scalability and efficiency with high-dimensional function approximators.

We will focus on the gradient-based derivation that underlies many policy optimization algorithms.

# Derivation of the Policy Gradient Theorem (I)

- Define the performance objective:

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} [G(\tau)] = \int_{\tau} P(\tau; \theta) G(\tau) d\tau,$$

where  $\tau = (s_0, a_0, s_1, a_1, \dots)$  denotes a trajectory, and

$$P(\tau; \theta) = \rho_0(s_0) \prod_{t=0}^{T-1} \pi_{\theta}(a_t \mid s_t) P(s_{t+1} \mid s_t, a_t).$$

- Our goal is to compute the gradient:

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \int_{\tau} P(\tau; \theta) G(\tau) d\tau.$$

# Derivation of the Policy Gradient Theorem (II)

- We apply the likelihood ratio trick. Recognize that

$$\nabla_{\theta} P(\tau; \theta) = P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta).$$

- Noting that  $P(\tau; \theta)$  factors through the policy, we have:

$$\log P(\tau; \theta) = \log \rho_0(s_0) + \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t | s_t) + \sum_{t=0}^{T-1} \log P(s_{t+1} | s_t, a_t).$$

- Since the environment dynamics  $P(s_{t+1} | s_t, a_t)$  and initial state distribution  $\rho_0(s_0)$  do not depend on  $\theta$ , their gradients vanish.

Therefore,

$$\nabla_{\theta} \log P(\tau; \theta) = \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t).$$

- Plugging this back into our gradient, we obtain:

$$\nabla_{\theta} J(\theta) = \int_{\tau} P(\tau; \theta) \left( \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) G(\tau) d\tau.$$

# Derivation of the Policy Gradient Theorem (III)

- Express the gradient as an expectation:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \right].$$

- In practice, it is common to use the *return from time  $t$*  rather than  $G(\tau)$  for each time step:

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

Then we write:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right].$$



# Introduction to REINFORCE

- **REINFORCE** is a Monte Carlo policy gradient algorithm.
- It directly uses complete episodes to estimate returns.
- The update rule is derived from the policy gradient theorem:

$$\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t \mid s_t) G_t \right],$$

where

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

- In words, we increase the likelihood of actions that yield higher returns.

# REINFORCE: Sampling and Return Computation

- 1 **Sampling a Trajectory:** Execute the current policy  $\pi_\theta$  to generate an episode:

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T).$$

This trajectory is sampled according to  $P(\tau; \theta)$ , as in the gradient theorem.

- 2 **Computing the Return:** For each time step  $t$ , calculate the return:

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

**Link to Theorem:** These  $G_t$  are the reward signals weighting the log-probability gradients in  $\nabla_\theta \log \pi_\theta(a_t | s_t) G_t$ .

# REINFORCE: Gradient Estimation & Policy Update

## 1 Gradient Estimation:

- For each time step  $t$  in the sampled trajectory, compute:

$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t).$$

- Multiply by the corresponding return  $G_t$  to form:

$$g_t = \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t.$$

- Recall:** This directly implements the term from the policy gradient theorem.

## 2 Policy Update: Aggregate the gradients over the episode and update parameters with a learning rate $\alpha$ :

$$\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-1} g_t.$$

## 3 Iteration: Repeat by collecting new trajectories, ensuring the policy continually improves.

# Variance Reduction & Baselines in REINFORCE

- **Problem:** The REINFORCE update

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t$$

can have high variance, slowing down learning.

- **Solution: Baselines** We subtract a baseline  $b(s_t)$  from the return:

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (G_t - b(s_t)).$$

- **Why it helps:**  $(G_t - b(s_t))$  represents the *advantage* of taking action  $a_t$  in state  $s_t$ . Using a baseline reduces the variance of the gradient estimate without introducing bias.
- **Common Baseline:** A natural choice for  $b(s_t)$  is the state-value function  $V^{\pi}(s_t)$ .

# Policy Methods: A Distinction

## Actor-Only Methods

- Learn the policy directly from experience.
- Do not estimate any value function.
- Example: **REINFORCE**.
- Simple to implement, but gradients have **high variance**.

## Actor-Critic Methods

- Combine a policy (Actor) with a value function (Critic).
- The Critic helps reduce variance by estimating how good actions are.
- Example: **PPO (Proximal Policy optimization)**.
- More stable and sample-efficient than Actor-Only.

# Actor-Critic Methods: Detailed Overview

- **Actor:**

- Parameterizes the policy as  $\pi_{\theta}(a \mid s)$  — it decides which action to take.
- Updated by ascending the gradient of the expected return.

- **Critic:**

- Estimates the state-value function  $V_w(s)$ , which approximates the expected return from state  $s$ .
- Provides feedback to the actor by evaluating the quality of its actions.

- **Interaction:**

- The Critic's evaluation is used as a *baseline* to compute the **advantage**:

$$A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s).$$

- In practice, a sample-based form is often used:

$$A_t = G_t - V_w(s_t),$$

where  $G_t$  is the return computed from time step  $t$ .

- **Benefit:** By using the critic as a baseline, we reduce the variance of the policy gradient estimate.

# Deriving the Actor's Gradient (Part I)

- **Performance Objective:**

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}}[G(\tau)]$$

where  $\tau$  is a trajectory generated by the policy.

- **Policy Gradient Theorem:**

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right],$$

with

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

- This expression indicates that we adjust  $\theta$  to increase the probability of actions leading to higher returns.

# Deriving the Actor's Gradient (Part II)

- **Introducing a Baseline:** Subtracting a baseline, such as the state-value function  $V_w(s_t)$ , does not change the expected value but reduces variance:

$$A_t = G_t - V_w(s_t).$$

- **Final Actor Update:** The gradient estimate becomes:

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A_t \right],$$

leading to the parameter update:

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta).$$

- This update reinforces actions whose outcomes surpass the expected value.



# Critic Update – TD Error and Loss (Part I)

- **Objective for the Critic:** Learn an approximation  $V_w(s)$  to the expected return from state  $s$ .
- **Temporal Difference (TD) Error:**

$$\delta_t = r_t + \gamma V_w(s_{t+1}) - V_w(s_t).$$

- $\delta_t$  measures the discrepancy between the current value estimate and a bootstrap estimate from the next state.

# Critic Update – Loss Minimization (Part II)

- **Loss Function:** To train the critic, we minimize the mean squared error (MSE) of the TD error:

$$L(w) = \frac{1}{2} \delta_t^2 = \frac{1}{2} \left( r_t + \gamma V_w(s_{t+1}) - V_w(s_t) \right)^2.$$

- **Parameter Update:** Update the critic parameters via gradient descent:

$$w \leftarrow w - \alpha_w \nabla_w L(w).$$

- A well-trained critic provides accurate value estimates, leading to reliable advantage computations.

# Actor-Critic Process (Part I): Data Collection & Advantage Estimation

- 1 **Sampling an Episode:** Execute the current policy  $\pi_\theta$  to obtain a trajectory:

$$\tau = \{s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T\}.$$

- 2 **Compute Returns:** For each time step  $t$ :

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

- 3 **Estimate Advantage:** Use the critic to compute:

$$A_t = G_t - V_w(s_t).$$

- 4 This estimation bridges the actor's policy update with the critic's evaluation.

# Actor-Critic Process (Part II): Updates and Iteration

- 1 **Update the Critic:** Minimize the loss:

$$L(w) = \frac{1}{2} \left( r_t + \gamma V_w(s_{t+1}) - V_w(s_t) \right)^2,$$

updating  $w$  via gradient descent.

- 2 **Update the Actor:** Adjust  $\theta$  using:

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$$

- 3 **Iteration:** Collect new episodes and repeat the process to continually refine both the policy and value estimates.

- **Instability in Naive Policy Gradients:**

- Large, unconstrained policy updates can lead to erratic and unstable learning.
- Sudden shifts in policy often degrade performance in complex environments.

- **Trust Region Policy Optimization (TRPO):**

- Introduces a KL-divergence constraint:

$$D_{\text{KL}}(\pi_{\theta}, \pi_{\theta_{\text{old}}}) \leq \delta.$$

- Limits policy updates to remain within a “trust region.”

# PPO Motivation (Part II)

- **Practicality of PPO:**

- PPO aims to achieve the stability of TRPO without the need for complex second-order optimization.

- **Key Idea:** Use a *clipped* objective to restrict the extent of policy updates.

- This prevents large deviations from the old policy without explicitly enforcing a KL constraint.

- **Outcome:** More stable updates, improved sample efficiency, and ease of implementation.

# From Standard Actor-Critic to PPO

- **Standard Actor-Critic Update:**

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \hat{A}_t.$$

- **Limitation:** Direct updates can lead to large changes in the policy if not controlled.
- **Introducing the Probability Ratio:**

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}.$$

- **Interpretation:**

- $r_t(\theta)$  measures the change in probability for the taken action between the new and old policies.
- Ideally,  $r_t(\theta)$  should remain near 1 to ensure moderate updates.

# PPO Clipped Objective

## Clipped Surrogate Objective

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[ \min \left( r_t(\theta) \hat{A}_t, \text{clip} \left( r_t(\theta), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right) \right].$$

- $\hat{A}_t$  is the advantage estimate.
- The objective takes the minimum of the unclipped and clipped terms to restrict large policy updates.
- **Role of Clipping:**
  - Prevents  $r_t(\theta)$  from straying far from 1 (i.e., a small change from the old policy).
  - If  $r_t(\theta)$  exceeds the range  $[1 - \epsilon, 1 + \epsilon]$ , the corresponding term is clipped.
- The hyperparameter  $\epsilon$  controls how much the new policy is allowed to deviate from the old policy during each update.
- **Outcome:** The update remains *proximal* to the prior policy, ensuring stable learning.



# Actor and Critic Loss in PPO (Part I)

## Actor Loss:

$$L_{\text{actor}}(\theta) \approx \mathbb{E}_t \left[ \min \left( r_t(\theta) \hat{A}_t, \text{clip} \left( r_t(\theta), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right) \right].$$

- This objective encourages policies that improve advantage estimates while avoiding excessive updates.

## Critic Loss:

$$L_{\text{critic}}(w) \approx \mathbb{E}_t \left[ \left( V_w(s_t) - G_t^{\text{target}} \right)^2 \right],$$

where  $G_t^{\text{target}}$  is a return estimate (e.g., one-step or n-step bootstrapped return).

# Actor and Critic Loss in PPO (Part II)

- **Combined Objective:**

$$L^{\text{PPO}}(\theta, w) = L_{\text{actor}}(\theta) - c_1 L_{\text{critic}}(w) + c_2 H(\pi_\theta),$$

where:

- $c_1$  is the coefficient balancing the critic loss.
- $c_2$  is the coefficient for the entropy bonus.
- **Entropy Bonus  $H(\pi_\theta)$ :**
  - Encourages exploration by penalizing overly deterministic policies.
- **Purpose:** Jointly optimizes the policy and value function, ensuring both accurate value predictions and stable policy improvement.

# Generalized Advantage Estimation (GAE): Fundamentals

- **Motivation:**

- Naively using  $G_t - V(s_t)$  yields a high-variance advantage estimate.
- A reliable advantage estimate is crucial for stable policy updates.

- **Temporal Difference (TD) Residual:**

$$\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t).$$

- $\delta_t$  captures the immediate error of our value prediction.

- **GAE Core Idea:**

- Use a weighted sum of future TD residuals to compute the advantage.

# Generalized Advantage Estimation (GAE): Computation and Benefits

- **GAE Formula:**

$$\hat{A}_t = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l},$$

where  $\lambda \in [0, 1]$  is a hyperparameter that controls the trade-off:

- $\lambda = 0$ : Uses only the immediate TD error (high bias, low variance).
  - $\lambda = 1$ : Approximates the full Monte Carlo return (low bias, high variance).
- **Benefits:**
    - **Variance Reduction:** Bootstrapping from  $V(s)$  reduces fluctuations.
    - **Timely Updates:** Utilizes multi-step information without waiting for episode termination.
    - **Tunable Bias-Variance Trade-Off:**  $\lambda$  can be adjusted to meet the needs of the environment.

# Full PPO Algorithm: Data Collection & Advantage Estimation

- **Step 1: Data Collection**

- Run the current policy  $\theta_{\text{old}}$  to collect a batch of trajectories.
- Gather tuples  $(s_t, a_t, r_t, s_{t+1})$  for each time step.

- **Step 2: Compute Returns**

- For each time step, calculate the return:

$$G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k.$$

- Returns summarize the future rewards from the current state, but using full returns can introduce high variance.

- **Step 3: Advantage Estimation**

- Estimate the advantages  $\hat{A}_t$  using methods like GAE.

# Full PPO Algorithm: Policy & Critic Updates

## • Step 4: Policy Update

- Compute the probability ratio:

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}.$$

- Optimize the clipped objective:

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[ \min \left( r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right].$$

## • Step 5: Critic Update

- Update the critic parameters  $w$  by minimizing:

$$L_{\text{critic}}(w) = \mathbb{E}_t \left[ (V_w(s_t) - G_t^{\text{target}})^2 \right].$$

- $G_t^{\text{target}}$  can use bootstrapping (e.g., one-step or multi-step returns) to further stabilize learning.

## • Final Step: Iteration

- After updating, set  $\theta_{\text{old}} \leftarrow \theta$  and repeat for the next data batch.

# Benefits of PPO

- **Stable Updates:** Clipping keeps policy changes moderate and avoids catastrophic updates.
- **Sample Efficiency:** Multiple epochs over the same batch improve the efficiency of data usage.
- **Simplicity:** Avoids the complex second-order computations required in TRPO.
- **Wide Applicability:** Effective in continuous control and high-dimensional RL tasks.

# Appendix

Additional Material and Derivations



# Derivation of the Likelihood Ratio Trick (Part I)

## Objective:

$$J(\theta) = \mathbb{E}_{x \sim p_\theta}[f(x)] = \int p_\theta(x) f(x) dx.$$

## Step 1: Compute the Gradient of $J(\theta)$

$$\nabla_\theta J(\theta) = \nabla_\theta \int p_\theta(x) f(x) dx = \int \nabla_\theta p_\theta(x) f(x) dx.$$

## Step 2: Log-Derivative Identity

$$\nabla_\theta \log p_\theta(x) = \frac{1}{p_\theta(x)} \nabla_\theta p_\theta(x) \quad \Rightarrow \quad \nabla_\theta p_\theta(x) = p_\theta(x) \nabla_\theta \log p_\theta(x).$$

# Derivation of the Likelihood Ratio Trick (Part II)

## Step 3: Substitute Back into the Gradient Expression

$$\nabla_{\theta} J(\theta) = \int p_{\theta}(x) \nabla_{\theta} \log p_{\theta}(x) f(x) dx.$$

## Step 4: Express as an Expectation

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{x \sim p_{\theta}} [f(x) \nabla_{\theta} \log p_{\theta}(x)].$$

**Intuition:** We rewrite the derivative of the probability  $p_{\theta}(x)$  as a product of  $p_{\theta}(x)$  and the gradient of its log — allowing us to express the full gradient as an expectation, which can be estimated from samples.